

AI Planning for Autonomy

4. Generating Heuristic Functions

How to Relax: Formally, and Informally, and During Search

Tim Miller and Nir Lipovetzky



THE UNIVERSITY OF
MELBOURNE

Winter Term 2019

Agenda

- 1 Motivation
- 2 How to Relax Informally
- 3 How to Relax Formally
- 4 How to Relax During Search
- 5 Conclusion

Motivation

→ “Relax”ing is a methodology to construct heuristic functions.

- You can use it when programming a solution to some problem you want/need to solve.
- Planning systems can use it to derive a heuristic function **automatically** from the planning task description (the PDDL input).
 - **Note 1:** If the user had to supply the heuristic function by hand, then we would lose our two main selling points (generality & autonomy & flexibility & rapid prototyping, cf. → **Lecture 1-2**).
 - **Note 2:** It can of course be of advantage to give the user the *possibility* to (conveniently) supply additional heuristics. Not covered in this course.

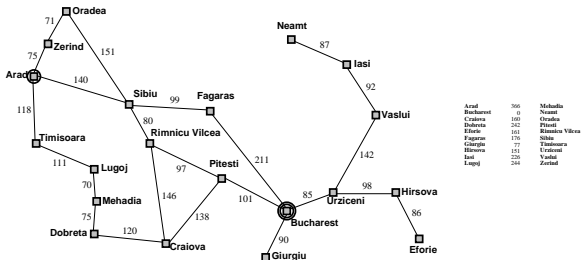
How to Relax Informally

How To Relax:

- You have a problem, \mathcal{P} , whose perfect heuristic h^* you wish to estimate.
- You define a **simpler problem**, \mathcal{P}' , whose perfect heuristic h'^* can be used to **estimate h^*** .
- You define a transformation, r , that **simplifies** instances from \mathcal{P} into instances \mathcal{P}' .
- Given $\Pi \in \mathcal{P}$, you estimate $h^*(\Pi)$ by $h'^*(r(\Pi))$.

→ Relaxation means to simplify the problem, and take the solution to the simpler problem as the heuristic estimate for the solution to the actual problem.

Relaxation in Route-Finding



How to derive straight-line distance by relaxation?

- Problem \mathcal{P} : Route finding.
- Simpler problem \mathcal{P}' :
- Perfect heuristic h'^* for \mathcal{P}' :
- Transformation r :

Relaxation in the 8-Puzzle

7	2	4
5		6
8	3	1

Start State

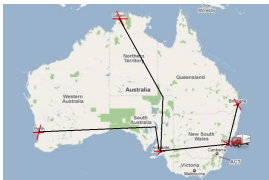
	1	2
3	4	5
6	7	8

Goal State

Perfect heuristic h^* for \mathcal{P} : Actions = “A tile can move from square A to square B if A is adjacent to B and B is blank.”

- How to derive the Manhattan distance heuristic?
- How to derive the misplaced tiles heuristic?
- h'^* (resp. r) in both: optimal cost in \mathcal{P}' (resp. use different actions).
- Here: Manhattan distance = , misplaced tiles = .

“Goal-Counting” Relaxation in Australia



- **Propositions** P : $at(x)$ for $x \in \{Sy, Ad, Br, Pe, Da\}$; $v(x)$ for $x \in \{Sy, Ad, Br, Pe, Da\}$.
- **Actions** $a \in A$: $drive(x, y)$ where x, y have a road; $pre_a = \{at(x)\}$, $add_a = \{at(y), v(y)\}$, $del_a = \{at(x)\}$.
- **Initial state** I : $at(Sy), v(Sy)$.
- **Goal** G : $at(Sy), v(x)$ for all x .

Let's “act as if we could achieve each goal directly”:

- **Problem** \mathcal{P} : All STRIPS planning tasks.
- **Simpler problem** \mathcal{P}' : All STRIPS planning tasks with empty preconditions and deletes.
- **Perfect heuristic** h^* for \mathcal{P}' : Optimal plan cost ($= h^*$).
- **Transformation** r :
- Heuristic value here? .

→ Optimal STRIPS planning with empty preconditions and deletes is still **NP**-hard! (Reduction from MINIMUM COVER, of goal set by add lists.)

→ Need to **approximate** the perfect heuristic h^* for \mathcal{P}' . Hence **goal counting**: just approximate h^* by number-of-false-goals.

How to Relax Formally: Before We Begin

- The definition on the next slide is not to be found in any textbook, and not even in any paper.
- Methods generating heuristic functions differ widely, and it is quite difficult (impossible?) to make *one* definition capturing them *all* in a natural way.
- Nevertheless, a formal definition is useful to state precisely what are the relevant distinction lines in practice.
- The present definition does, I think, do a rather good job of this.
 - It nicely fits what is currently used in planning.
 - It is flexible in the distinction lines, and it captures the basic construction, as well as the essence of all relaxation ideas.

Relaxations


Definition (Relaxation). Let $h^* : \mathcal{P} \mapsto \mathbb{R}_0^+ \cup \{\infty\}$ be a function. A *relaxation* of h^* is a triple $\mathcal{R} = (\mathcal{P}', r, h'^*)$ where \mathcal{P}' is an arbitrary set, and $r : \mathcal{P} \mapsto \mathcal{P}'$ and $h'^* : \mathcal{P}' \mapsto \mathbb{R}_0^+ \cup \{\infty\}$ are functions so that, for all $\Pi \in \mathcal{P}$, the *relaxation heuristic* $h^{\mathcal{R}}(\Pi) := h'^*(r(\Pi))$ satisfies $h^{\mathcal{R}}(\Pi) \leq h^*(\Pi)$. The relaxation is:

- *native* if $\mathcal{P}' \subseteq \mathcal{P}$ and $h'^* = h^*$;
- *efficiently constructible* if there exists a polynomial-time algorithm that, given $\Pi \in \mathcal{P}$, computes $r(\Pi)$;
- *efficiently computable* if there exists a polynomial-time algorithm that, given $\Pi' \in \mathcal{P}'$, computes $h'^*(\Pi')$.

Reminder:

- You have a problem, \mathcal{P} , whose perfect heuristic h^* you wish to estimate.
- You define a simpler problem, \mathcal{P}' , whose perfect heuristic h'^* can be used to (admissibly!) estimate h^*
- You define a transformation, r , from \mathcal{P} into \mathcal{P}' .
- Given $\Pi \in \mathcal{P}$, you estimate $h^*(\Pi)$ by $h'^*(r(\Pi))$.


Relaxations: Illustration



Example route-finding:

- Problem \mathcal{P} : Route finding.
- Simpler problem \mathcal{P}' :
- Perfect heuristic h'^* for \mathcal{P}' :
- Transformation r :

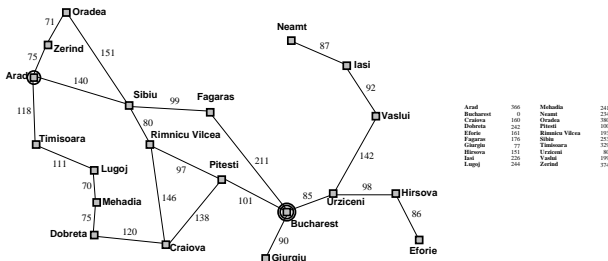
Native Relaxations: Illustration



Example “goal-counting”:

- Problem \mathcal{P} : All STRIPS planning tasks.
- Simpler problem \mathcal{P}' : All STRIPS planning tasks with empty preconditions and deletes.
- Perfect heuristic h'^* for \mathcal{P}' : Optimal plan cost = h^* .
- Transformation r :

Relaxation in Route-Finding: Properties



Relaxation $\mathcal{R} = (\mathcal{P}', r, h'^*)$: Pretend you're a bird.

- Native?
- Efficiently constructible?
- Efficiently computable?

Relaxation in the 8-Puzzle: Properties

7	2	4
5		6
8	3	1

Start State

	1	2
3	4	5
6	7	8

Goal State

Relaxation $\mathcal{R} = (\mathcal{P}', r, h'^*)$: Use more generous actions rule to obtain Manhattan distance.

- Native?
- Efficiently constructible?
- Efficiently computable?

What shall we do with the relaxation?

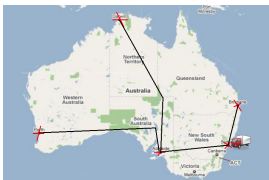
What if \mathcal{R} is not efficiently constructible?

- Either (a) approximate r , or (b) design r in a way so that it will typically be feasible, or (c) just live with it and hope for the best.
- Vast majority of known relaxations (in planning) are efficiently constructible.

What if \mathcal{R} is not efficiently computable?

- Either (a) approximate h'^* , or (b) design h'^* in a way so that it will typically be feasible, or (c) just live with it and hope for the best.
- Many known relaxations (in planning) are efficiently computable, some aren't. The latter use (a); (b) and (c) are not used anywhere right now.

“Goal-Counting” Relaxation in Australia: Properties



- **Propositions** P : $at(x)$ for $x \in \{Sy, Ad, Br, Pe, Da\}$; $v(x)$ for $x \in \{Sy, Ad, Br, Pe, Da\}$.
- **Actions** $a \in A$: $drive(x, y)$ where x, y have a road; $pre_a = \{at(x)\}$, $add_a = \{at(y), v(y)\}$, $del_a = \{at(x)\}$.
- **Initial state** I : $at(Sy), v(Sy)$.
- **Goal** G : $at(Sy), v(x)$ for all x .

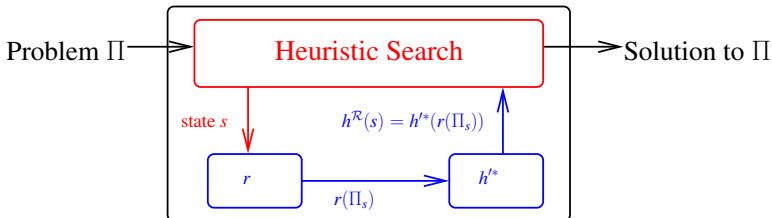
Relaxation $\mathcal{R} = (\mathcal{P}', r, h'^*)$: Remove preconditions and deletes, then use h^* .

- Native?
- Efficiently constructible?
- Efficiently computable?

What shall we do with the relaxation?

How to Relax During Search: Diagram

Using a relaxation $\mathcal{R} = (\mathcal{P}', r, h'^*)$ during search:



→ Π_s : Π with initial state replaced by s , i.e., $\Pi = (F, A, c, I, G)$ changed to (F, A, c, s, G) .

→ The task of finding a plan for search state s .

→ We will be using this notation in the course!

Questionnaire

Question!

Say we have a robot with one gripper, two rooms A and B , and n balls we must transport. The actions available are $move_{XY}$, $pick_B$ and $drop_B$; say h = "number of balls not yet in room B ". Can h be derived as $h^{\mathcal{R}}$ for a relaxation \mathcal{R} ?

(A): No.

(B): Yes, just drop the deletes

(C): Sure, every admissible h can be derived via a relaxation.

(D): I'd rather relax at the beach.

Summary

- **Relaxation** is a method to compute heuristic functions.
- Given a problem \mathcal{P} we want to solve, we define a **relaxed problem** \mathcal{P}' . We derive the heuristic by mapping into \mathcal{P}' and taking the solution to this simpler problem as the heuristic estimate.
- Relaxations can be **native**, **efficiently constructible**, and/or **efficiently computable**. None of this is a strict requirement to be useful.
- During search, **the relaxation is used *only inside* the computation of the heuristic function on each state**; the relaxation does not affect anything else. (This can be a bit confusing especially for native relaxations like ignoring deletes.)

Remarks

The goal-counting approximation h = “count the number of goals currently not true” is **a very uninformative heuristic function**:

- Range of heuristic values is small ($0 \dots |G|$).
- We can transform any planning task into an equivalent one where $h(s) = 1$ for all non-goal states s . **How?**
- Ignores almost all structure: Heuristic value does not depend on the actions at all!

→ By the way, is h safe/goal-aware/admissible/consistent?

→ We will see in → **the next lecture** how to compute **much** better heuristic functions.